

Vikram Kumar Parmar

Data Engineer | Austin, TX

585-309-8998 | parmar.vik25@gmail.com | github.com/markiv25 | linkedin.com/in/vikramparmar25

PROFESSIONAL SUMMARY

Data Engineer running a real-time analytics platform that processes 500M events/day for aviation clients on Kafka, Airflow, and ClickHouse at Anuvu. Led a 100TB migration from MariaDB to ClickHouse and own the Airflow and Kafka platforms for a 4-engineer data team. Drive AI workflow adoption across engineering, building internal Claude-based automations that replaced a 48-hour manual reporting cycle with a sub-4-hour workflow across 8 airline accounts.

SKILLS

Languages: Python, SQL, Bash

Data Engineering: Apache Kafka, Apache Airflow, Apache Spark, PySpark, Polars, Parquet, CDC, ETL/ELT, dimensional modeling, data contracts, schema enforcement, lineage tracking

Databases: ClickHouse, MariaDB, MySQL, SQL Server

Cloud and Infrastructure: AWS (S3, RDS, EC2, SQS), Redis, Linux, Git, CI/CD

Observability and BI: Grafana, Tableau

AI and LLM Tooling: Claude API, agentic workflow design, internal Claude-based automation tooling, LLM-assisted data validation

TECHNICAL EXPERIENCE

Anuvu | Data Engineer, Streaming Platform

January 2023 – Present | Remote

- Designed and built the Kafka-based CDC pipeline that powers a real-time analytics platform processing 500M events/day for aviation clients on Kafka, Airflow, and ClickHouse; took data freshness from hourly batch to sub-minute, powering operational dashboards, customer-facing analytics, and internal alerting downstream.
- Led a 3-person team on a 100TB, 200-table migration from MariaDB to ClickHouse; cut analytics dashboard latency from 45s to under 5s, improved write throughput 3x, and cut MariaDB infrastructure spend by ~\$23K/month post-decommissioning.
- Redesigned the per-flight processing path from an SQS + RDS pipeline to an in-process Polars and DuckDB pipeline; cut per-flight runtime from 2–2.5 min to 1.5–1.7 min (25–30% faster) and reduced load on the shared RDS instance by eliminating the round-trip for every flight.
- Platform owner for Apache Airflow and Kafka in the data engineering team; run 10 production DAGs processing 1–5 TB/day of telemetry on AWS (S3, RDS, EC2, SQS, Redis), with cluster upgrades, on-call ownership, and DAG governance shared across the team.
- Modernized a Python 2 batch parser to a Python 3 PySpark/Polars pipeline, delivering ~1.5x faster end-to-end runtime versus the legacy baseline on daily telemetry and unblocking reprocessing windows that were previously infeasible inside the SLA window.
- Built a data quality and contracts framework (schema validation, lineage tracking, alerting) integrated into Airflow; cut ad-hoc stakeholder data requests by ~35% (measured against prior-quarter ticket volume) by shifting routine reporting onto governed self-serve dashboards.
- Drive AI adoption across a 32-person engineering team: cut client reporting turnaround from 48 hours to under 4 hours across 8 airline accounts via an agent-driven workflow, and built an internal Claude-based toolkit (7 skills and agents for pipeline scaffolding, data validation, and token-budget optimization) actively used by most of the engineering team.

Rochester Institute of Technology | Data Engineer Intern

August 2022 – December 2022 | Rochester, NY

- Built ETL workflows (Informatica, SQL Server, Azure) feeding a departmental data warehouse used for self-service analytics across 5 departments; cut data retrieval times ~30% and maintained pipelines covering 15,000 student records.

Ascension Health | Master Data Management Intern

May 2022 – August 2022 | Remote

- Designed Python ingestion pipelines processing 500K records/week; delivered KPI reporting across 3 departments and improved inventory forecast accuracy, reducing overstock ~15%.

EDUCATION

Rochester Institute of Technology – Rochester, NY

M.S., Information Technology and Analytics | 2020 – 2023

University of Mumbai – Mumbai, India

B.E., Electronics Engineering | 2015 – 2019

PROJECTS

slim-margin – Open-source Claude Code plugin enforcing token-budget discipline and Opus/Sonnet/Haiku tier routing on long-context tasks. MIT-licensed; github.com/markiv25/slim-margin

PUBLICATIONS

Srinivasan, A., Parmar, V., et al. (2021). *Anomaly Detection System for Smart Home using Machine Learning*. IEEE ICSSA, pp. 52–55. DOI: 10.1109/ICSSA53632.2021.00018